

Introduction to R

PSY517 Quantitative Analysis III

Derek Powell

August 2, 2021

2021-08-24

└─ Why R? (or, why am I doing this to you?)

poll class

- how many, when you heard class would use R, felt excited?
 - how many felt worried?
- how many have used R? what have you done with it?
- How was the LSR reading?

Why R? (or, why am I doing this to you?)

- R is a powerful tool for statistics
- R is a powerful tool for data management
 - Scripts and automation reduce errors, reduce tedium, and make analyses reproducible
- Programming is where the \$\$\$ is at
 - Basic programming knowledge is becoming a necessity in industry and academia
- R has a great community, tons of support, and is really not that hard to learn
- We will substitute programming in place of math to understand statistics

└ Why R? (or, why am I doing this to you?)

- R is a powerful tool for statistics
- R is a powerful tool for data management
 - Scripts and automation reduce errors, reduce tedium, and make analyses reproducible
- Programming is where the \$\$\$ is at
 - Basic programming knowledge is becoming a necessity in industry and academia
- R has a great community, tons of support, and is really not that hard to learn
- We will substitute programming in place of math to understand statistics

2021-08-24

└ Kicking ass

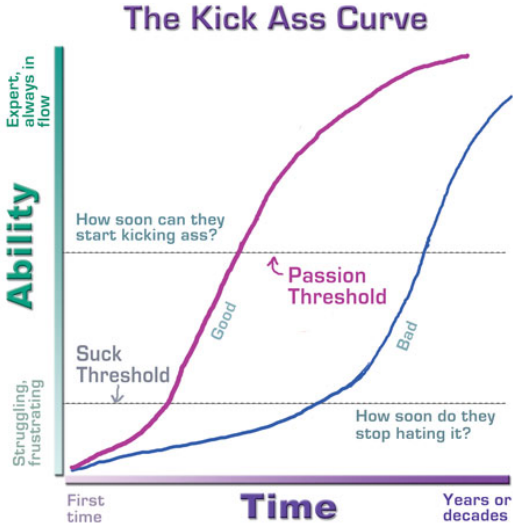


Figure 1: The Kick-Ass curve

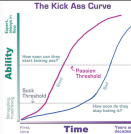


Figure 1: The Kick-Ass curve

Let's jump in

I'll show you ...

- scripts
- notebooks
- how to load data
- how to manipulate data
- how to plot data
- how to conduct a statistical test

Let's jump in

I'll show you ...

- scripts
- notebooks
- how to load data
- how to manipulate data
- how to plot data
- how to conduct a statistical test

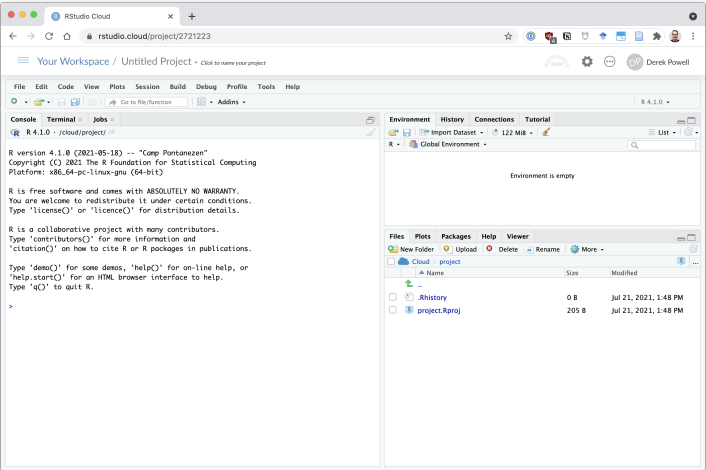


Figure 2: The Rstudio interface

2021-08-24

Rstudio tour

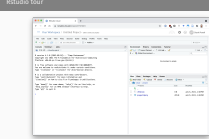


Figure 2: The Rstudio interface

Interactive coding demo

2021-08-24

Introduction to R
└ Interactive coding demo

Interactive coding demo

Let's jump over to Rstudio, open a script, and run some code.

2021-08-24

Introduction to R
└─ Interactive coding demo

└─ Interactive coding demo

Let's jump over to Rstudio, open a script, and run some code.

R “packages” are collections of code, functions, and data that have been packaged together. Typically, packages support conducting certain types of analyses, making plots, etc.

The most important set of packages we will learn in this course are the **tidyverse** packages. We can load all of them at once with the command below:

```
library(tidyverse)
```

2021-08-24

Introduction to R

└─ Interactive coding demo

└─ Loading packages

Loading packages

R “packages” are collections of code, functions, and data that have been packaged together. Typically, packages support conducting certain types of analyses, making plots, etc.

The most important set of packages we will learn in this course are the **tidyverse** packages. We can load all of them at once with the command below:

```
library(tidyverse)
```

```
df <- read_csv("stroop-2014.csv")
```

- Load data from Daniel Lakens' introductory psychology course.
- Students performed the classic "Stroop" task: name ink color that congruent or incongruent words were printed in
 - (e.g. The word "Green" printed in green versus yellow ink).
- Reading interferes with color naming, so it takes longer to name the ink color when the word it is printing is "incongruent"

2021-08-24

Introduction to R
└─ Interactive coding demo

└─ Loading data

```
df <- read_csv("stroop-2014.csv")
```

- Load data from Daniel Lakens' introductory psychology course.
- Students performed the classic "Stroop" task: name ink color that congruent or incongruent words were printed in
 - (e.g. The word "Green" printed in green versus yellow ink).
- Reading interferes with color naming, so it takes longer to name the ink color when the word it is printing is "incongruent"

Inspecting data with head()

We can inspect our data with the `head()` function.

```
head(df)
```

```
## # A tibble: 6 x 3
##   subj_num congruent incongruent
##   <dbl>      <dbl>      <dbl>
## 1         1    13.7        22.7
## 2         2    14.8        25.9
## 3         3    16.8        25.7
## 4         4    11.9        21.2
## 5         5    10.5        22.6
## 6         6     9.44        20.7
```

Introduction to R

Interactive coding demo

Inspecting data with head()

We can inspect our data with the `head()` function.

```
head(df)

## # A tibble: 6 x 3
##   subj_num congruent incongruent
##   <dbl>      <dbl>      <dbl>
## 1         1    13.7        22.7
## 2         2    14.8        25.9
## 3         3    16.8        25.7
## 4         4    11.9        21.2
## 5         5    10.5        22.6
## 6         6     9.44        20.7
```

2021-08-24

Our `df` object is a dataframe that consists of multiple variables. We can “index” or extract those variables using the `$` operator. Let’s pull out the congruent trial data.

```
df$congruent
```

```
## [1] 13.741 14.788 16.819 11.888 10.516  9.436 13.256 18.643 15.827 13.000
## [11] 17.600 19.065 16.387 19.765 10.281 16.702 17.465 19.040 22.678 15.767
## [21] 14.445 16.888 10.034 10.091 16.460 15.721 10.900 17.196 21.392 17.725
## [31] 23.204 15.926 16.618 14.850 13.003 15.247  9.396 11.998 15.615 22.891
## [41] 10.825 14.690 21.000 18.363  9.733 14.563 12.162 13.076
```

Our `df` object is a dataframe that consists of multiple variables. We can “index” or extract those variables using the `$` operator. Let’s pull out the congruent trial data.

```
df$congruent
```

```
## [1] 13.741 14.788 16.819 11.888 10.516  9.436 13.256 18.643 15.827 13.000
## [11] 17.600 19.065 16.387 19.765 10.281 16.702 17.465 19.040 22.678 15.767
## [21] 14.445 16.888 10.034 10.091 16.460 15.721 10.900 17.196 21.392 17.725
## [31] 23.204 15.926 16.618 14.850 13.003 15.247  9.396 11.998 15.615 22.891
## [41] 10.825 14.690 21.000 18.363  9.733 14.563 12.162 13.076
```

We can compute the mean and standard deviation of our congruent trials with:

```
mean(df$congruent)
```

```
## [1] 15.34742
```

```
sd(df$congruent)
```

```
## [1] 3.700337
```

Note that `mean()` and `sd()` are *functions* that perform some computation on whatever inputs they are given.

R includes many built-in functions, but sometimes we would like to do something new and unique. In those cases, we can define our own functions.

2021-08-24

Introduction to R

└ Interactive coding demo

└ Using and defining functions

Using and defining functions

```
We can compute the mean and standard deviation of our congruent trials with:
mean(df$congruent)

## [1] 15.34742
sd(df$congruent)

## [1] 3.700337

Note that mean() and sd() are functions that perform some computation on
whatever inputs they are given.

R includes many built-in functions, but sometimes we would like to do something
new and unique. In those cases, we can define our own functions.
```

One major omission in R is a built-in function to compute the standard error of a variable. R has a function `sd()` to compute the standard deviation, but nothing for the standard error. Recall, the formula for the standard error is:

$$SE = \frac{\sigma}{\sqrt{n}}$$

Let's make our own standard error function in R, named `se()`.

```
se <- function(x){  
  sd(x) / sqrt(length(x))  
}
```

To learn more about this function, try typing `?sqrt` and `?length` into the R console.

2021-08-24

Introduction to R

└ Interactive coding demo

└ Defining functions

Try computing each of the components of this function, `sd`, `sqrt`, and `length`

Defining functions

One major omission in R is a built-in function to compute the standard error of a variable. R has a function `sd()` to compute the standard deviation, but nothing for the standard error. Recall, the formula for the standard error is:

$$SE = \frac{\sigma}{\sqrt{n}}$$

Let's make our own standard error function in R, named `se()`.

```
se <- function(x){  
  sd(x) / sqrt(length(x))  
}
```

To learn more about this function, try typing `?sqrt` and `?length` into the R console.

Now we can use our function to calculate the standard error.

```
Checking each piece  
sd(df$congruent)
```

```
## [1] 3.700337
```

```
length(df$congruent)
```

```
## [1] 48
```

```
sqrt(48)
```

```
## [1] 6.928203
```

```
3.700337 / 6.928203
```

```
## [1] 0.5340977
```

Using our function

```
se(df$congruent)
```

```
## [1] 0.5340977
```

Now we can use our function to calculate the standard error

```
Checking each piece  
sd(df$congruent)  
## [1] 3.700337  
length(df$congruent)  
## [1] 48  
sqrt(48)  
## [1] 6.928203  
3.700337 / 6.928203  
## [1] 0.5340977
```

Hands-on: Making your own function

Let's make a function to standardize the **congruent** and **incongruent** variables.

Recall our standard error equation and corresponding R function.

$$SE = \frac{\sigma}{\sqrt{n}}$$

```
se <- function(x){  
  sd(x) / sqrt(length(x))  
}
```

How would you create a function to standardize a variable based on the equation?

$$Z = \frac{x - \mu}{\sigma}$$

Introduction to R

Interactive coding demo

Hands-on: Making your own function

2021-08-24

Hands-on: Making your own function

Let's make a function to standardize the **congruent** and **incongruent** variables.

Recall our standard error equation and corresponding R function.

$$SE = \frac{\sigma}{\sqrt{n}}$$

```
se <- function(x){  
  sd(x) / sqrt(length(x))  
}
```

How would you create a function to standardize a variable based on the equation?

$$Z = \frac{x - \mu}{\sigma}$$

Working with data

2021-08-24

Introduction to R
└ Working with data

Working with data

First let's check if we are missing any values.

```
any(is.na(df))
```

```
## [1] FALSE
```

It looks like we aren't missing any data.

1. I passed our `df` object to the `is.na()` function, which will gave a matrix of `TRUE/FALSE` values indicating whether the value was missing or not.
2. Then I passed the output of `is.na()` into the `any()` function, which tells us if any of the values we give it are `TRUE`.

2021-08-24

Introduction to R
└ Working with data

└ Checking for missing values

Checking for missing values

First let's check if we are missing any values.

```
any(is.na(df))
```

```
## [1] FALSE
```

It looks like we aren't missing any data.

1. I passed our `df` object to the `is.na()` function, which will give a matrix of `TRUE/FALSE` values indicating whether the value was missing or not.
2. Then I passed the output of `is.na()` into the `any()` function, which tells us if any of the values we give it are `TRUE`.

Removing cases with `filter()`

Suppose participant #3 called us up with a confession: they were really distracted the day of the experiment and could hardly pay any attention. Let's remove them from our dataset.

We will use the “pipe” operator `%>%` to pipe our data (`df`) to the `filter()` function. We can compare two variables with `==`, `!=`, `>`, `<`, `>=`, and `<=`.

```
df %>%  
  filter(subj_num != 3)
```

```
## # A tibble: 47 x 3  
##   subj_num congruent incongruent  
##   <dbl>     <dbl>     <dbl>  
## 1         1      13.7      22.7  
## 2         2      14.8      25.9  
## 3         4      11.9      21.2  
## 4         5      10.5      22.6  
## # ... with 43 more rows
```

Introduction to R

Working with data

Removing cases with `filter()`

2021-08-24

Removing cases with `filter()`

Suppose participant #3 called us up with a confession: they were really distracted the day of the experiment and could hardly pay any attention. Let's remove them from our dataset.

We will use the “pipe” operator `%>%` to pipe our data (`df`) to the `filter()` function. We can compare two variables with `==`, `!=`, `>`, `<`, `>=`, and `<=`.

```
df %>%  
  filter(subj_num != 3)
```

```
## # A tibble: 47 x 3  
##   subj_num congruent incongruent  
##   <dbl>     <dbl>     <dbl>  
## 1         1      13.7      22.7  
## 2         2      14.8      25.9  
## 3         4      11.9      21.2  
## 4         5      10.5      22.6  
## # ... with 43 more rows
```

Creating new variables with mutate()

Things might feel more scientific if we store our response time measures as milliseconds (such precision!). We can compute new variables with the `mutate()` function.

```
df %>%  
  mutate(  
    congruent_ms = congruent*1000,  
    incongruent_ms = incongruent*1000  
  )
```

```
## # A tibble: 48 x 5  
##   subj_num congruent incongruent congruent_ms incongruent_ms  
##   <dbl>     <dbl>      <dbl>      <dbl>      <dbl>  
## 1         1    13.7        22.7    13741     22715  
## 2         2    14.8        25.9    14788     25916  
## 3         3    16.8        25.7    16819     25677  
## 4         4    11.9        21.2    11888     21213  
## 5         5    10.5        22.6    10516     22556  
## 6         6     9.44       20.7     9436     20715  
## # ... with 42 more rows
```

Introduction to R

Working with data

Creating new variables with mutate()

2021-08-24

Creating new variables with mutate()

Things might feel more scientific if we store our response time measures as milliseconds (such precision!). We can compute new variables with the `mutate()` function.

```
df %>%  
  mutate(  
    congruent_ms = congruent*1000,  
    incongruent_ms = incongruent*1000  
  )
```

```
## # A tibble: 48 x 5  
##   subj_num congruent incongruent congruent_ms incongruent_ms  
##   <dbl>     <dbl>      <dbl>      <dbl>      <dbl>  
## 1         1    13.7        22.7    13741     22715  
## 2         2    14.8        25.9    14788     25916  
## 3         3    16.8        25.7    16819     25677  
## 4         4    11.9        21.2    11888     21213  
## 5         5    10.5        22.6    10516     22556  
## 6         6     9.44       20.7     9436     20715  
## # ... with 42 more rows
```

Renaming variables with `rename()`

Maybe we'd like to capitalize our condition name variables.

```
df %>%  
  rename(Incongruent = incongruent, Congruent = congruent)
```

```
## # A tibble: 48 x 3  
##   subj_num Congruent Incongruent  
##   <dbl>     <dbl>     <dbl>  
## 1         1      13.7      22.7  
## 2         2      14.8      25.9  
## 3         3      16.8      25.7  
## 4         4      11.9      21.2  
## 5         5      10.5      22.6  
## 6         6       9.44      20.7  
## # ... with 42 more rows
```

Introduction to R

Working with data

Renaming variables with `rename()`

```
Renaming variables with rename()  
  
Maybe we'd like to capitalize our condition name variables.  
df %>%  
  rename(Incongruent = incongruent, Congruent = congruent)  
  
## # A tibble: 48 x 3  
##   subj_num Congruent Incongruent  
##   <dbl>     <dbl>     <dbl>  
## 1         1      13.7      22.7  
## 2         2      14.8      25.9  
## 3         3      16.8      25.7  
## 4         4      11.9      21.2  
## 5         5      10.5      22.6  
## 6         6       9.44      20.7  
## # ... with 42 more rows
```

2021-08-24

Stringing together actions with %>%

- The %>% operator pipes the output of one function into the next function as its first argument.

- For example,

```
any(is.na(df))
```

is the same as

```
df %>% is.na() %>% any()
```

- This lets us string commands together, as shown here

Example

```
df %>%  
  filter(subj_num != 3) %>%  
  mutate(  
    congruent_ms = congruent*1000,  
    incongruent_ms = incongruent*1000  
  ) %>%  
  rename(  
    Incongruent = incongruent,  
    Congruent = congruent  
  )
```

Introduction to R

Working with data

Stringing together actions with %>%

2021-08-24

- The %>% operator pipes the output of one function into the next function as its first argument.
- For example,

```
any(is.na(df))
```
- is the same as

```
df %>% is.na() %>% any()
```
- This lets us string commands together, as shown here

```
Example  
df %>%  
  filter(subj_num != 3) %>%  
  mutate(  
    congruent_ms = congruent*1000,  
    incongruent_ms = incongruent*1000  
  ) %>%  
  rename(  
    Incongruent = incongruent,  
    Congruent = congruent  
  )
```

Wide format

case	x	y
a	2	4
b	3	6

Long format

case	variable	value
a	x	2
a	y	4
b	x	2
b	y	6

The same data can be stored in different formats.

- **Wide format:** each row of a data table is a case with many variables.
- **Long format:** each row stores the value for one variable of one case.

2021-08-24

Tidying data

Wide format

case	x	y
a	2	4
b	3	6

The same data can be stored in different formats.

- **Wide format:** each row of a data table is a case with many variables.
- **Long format:** each row stores the value for one variable of one case.

Long format

case	variable	value
a	x	2
a	y	4
b	x	2
b	y	6

From wide to long with `gather()`

Let's convert our data from wide to long format using the `gather()` function.

```
df_long <- df %>%  
  gather(trial_type, rt, congruent, incongruent)
```

```
## # A tibble: 96 x 3  
##   subj_num trial_type    rt  
##   <dbl> <chr>      <dbl>  
## 1      1 congruent    13.7  
## 2      1 incongruent  22.7  
## 3      2 congruent    14.8  
## 4      2 incongruent  25.9  
## 5      3 congruent    16.8  
## 6      3 incongruent  25.7  
## # ... with 90 more rows
```

Introduction to R Working with data

From wide to long with `gather()`

2021-08-24

From wide to long with `gather()`

```
Let's convert our data from wide to long format using the gather() function.  
df_long <- df %>%  
  gather(trial_type, rt, congruent, incongruent)  
  
## # A tibble: 96 x 3  
##   subj_num trial_type    rt  
##   <dbl> <chr>      <dbl>  
## 1      1 congruent    13.7  
## 2      1 incongruent  22.7  
## 3      2 congruent    14.8  
## 4      2 incongruent  25.9  
## 5      3 congruent    16.8  
## 6      3 incongruent  25.7  
## # ... with 90 more rows
```


From long to wide with spread()

Sometimes wide data is what we want. We can convert long data to wide data with the `spread()` function.

```
df_wide <- df_long %>%  
  spread(trial_type, rt)
```

```
## # A tibble: 48 x 3  
##   subj_num congruent incongruent  
##   <dbl>      <dbl>      <dbl>  
## 1         1      13.7        22.7  
## 2         2      14.8        25.9  
## 3         3      16.8        25.7  
## 4         4      11.9        21.2  
## 5         5      10.5        22.6  
## 6         6       9.44        20.7  
## # ... with 42 more rows
```

Introduction to R

Working with data

From long to wide with spread()

2021-08-24

From long to wide with spread()

```
Sometimes wide data is what we want. We can convert long data to wide data with  
the spread() function.  
df_wide <- df_long %>%  
  spread(trial_type, rt)  
  
## # A tibble: 48 x 3  
##   subj_num congruent incongruent  
##   <dbl>      <dbl>      <dbl>  
## 1         1      13.7        22.7  
## 2         2      14.8        25.9  
## 3         3      16.8        25.7  
## 4         4      11.9        21.2  
## 5         5      10.5        22.6  
## 6         6       9.44        20.7  
## # ... with 42 more rows
```

A recent update to the **tidyverse** is the introduction of the **`pivot_longer()`** and **`pivot_wider()`** functions.

- **`pivot_longer()`** = fancier but more complicated **`gather()`**
- **`pivot_wider()`** = fancier but more complicated **`spread()`**

I don't expect we will benefit from the extra complexity of **`pivot_wider()`** and **`pivot_longer()`** for most of what we do in this class.

2021-08-24

Introduction to R
└─ Working with data

└─ Pivot functions

Pivot functions

A recent update to the **tidyverse** is the introduction of the **`pivot_longer()`** and **`pivot_wider()`** functions.

- **`pivot_longer()`** = fancier but more complicated **`gather()`**
- **`pivot_wider()`** = fancier but more complicated **`spread()`**

I don't expect we will benefit from the extra complexity of **`pivot_wider()`** and **`pivot_longer()`** for most of what we do in this class.

- `group_by()` breaks the data into groups
- `summarize()` can be used with functions that take a vector input (from each group) and output a single number

```
df_long %>%  
  group_by(trial_type) %>%  
  summarize(mean_rt = mean(rt), se_rt = se(rt))
```

```
## # A tibble: 2 x 3  
##   trial_type mean_rt se_rt  
##   <chr>      <dbl> <dbl>  
## 1 congruent    15.3 0.534  
## 2 incongruent  23.5 0.715
```

```
• group_by() breaks the data into groups  
• summarize() can be used with functions that take a vector input (from each group) and output a single number  
  
df_long %>%  
  group_by(trial_type) %>%  
  summarize(mean_rt = mean(rt), se_rt = se(rt))  
  
## # A tibble: 2 x 3  
##   trial_type mean_rt se_rt  
##   <chr>      <dbl> <dbl>  
## 1 congruent    15.3 0.534  
## 2 incongruent  23.5 0.715
```

Plotting data

2021-08-24

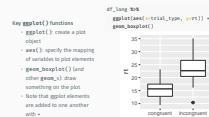
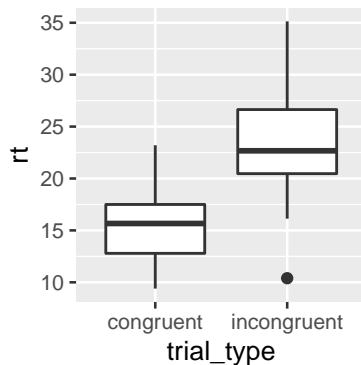
Introduction to R
└─ Plotting data

Plotting data

Key `ggplot()` functions

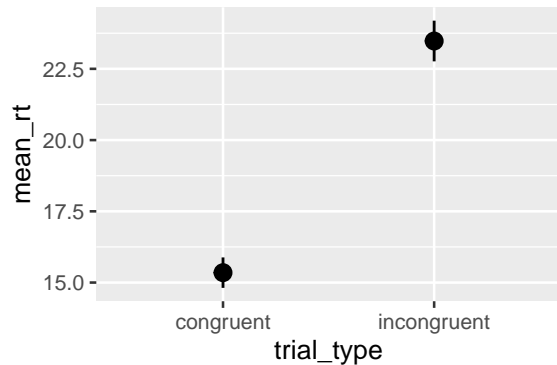
- `ggplot()`: create a plot object
- `aes()`: specify the mapping of variables to plot elements
- `geom_boxplot()` (and other `geom_s`): draw something on the plot
- Note that `ggplot` elements are added to one another with `+`

```
df_long %>%
  ggplot(aes(x=trial_type, y=rt)) +
  geom_boxplot()
```



Plotting summarized data

```
df_long %>%  
  group_by(trial_type) %>%  
  summarize(mean_rt = mean(rt), se_rt = se(rt)) %>%  
  mutate(ul = mean_rt + se_rt, ll = mean_rt - se_rt) %>%  
  ggplot(aes(x=trial_type, y = mean_rt, ymin=ll, ymax=ul)) +  
  geom_pointrange()
```

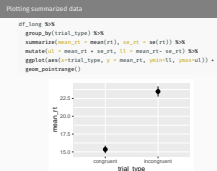


Introduction to R

Plotting data

Plotting summarized data

2021-08-24



Creating publication-quality figures (1)

```
df_long %>%
  group_by(trial_type) %>%
  summarize(
    mean_rt = mean(rt),
    se_rt = se(rt)
  ) %>%
  mutate(
    ul = mean_rt + 1.96 * se_rt,
    ll = mean_rt - 1.96 * se_rt
  ) %>%
  ggplot(
    aes(
      x = trial_type, y = mean_rt,
      ymin = ll, ymax = ul
    )
  ) +
  geom_bar(
    stat = "identity", width = .5, fill = "grey"
  ) +
  geom_errorbar(width = .1) +
  theme_bw(base_size = 28) +
  theme(panel.grid = element_blank()) +
  labs(
    x = "Trial Type",
    y = "Response time (s)"
  )
```

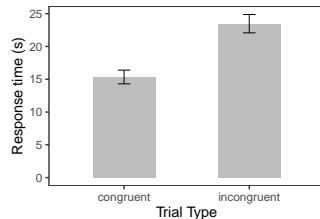


Figure 3: Mean response times by trial type with 95% confidence intervals

Introduction to R

Plotting data

Creating publication-quality figures (1)

2021-08-24

Creating publication-quality figures (1)

```
df_long %>%
  group_by(trial_type) %>%
  summarize(
    mean_rt = mean(rt),
    se_rt = se(rt)
  ) %>%
  mutate(
    ul = mean_rt + 1.96 * se_rt,
    ll = mean_rt - 1.96 * se_rt
  ) %>%
  ggplot(
    aes(
      x = trial_type, y = mean_rt,
      ymin = ll, ymax = ul
    )
  ) +
  geom_bar(
    stat = "identity", width = .5, fill = "grey"
  ) +
  geom_errorbar(width = .1) +
  theme_bw(base_size = 28) +
  theme(panel.grid = element_blank()) +
  labs(
    x = "Trial Type",
    y = "Response time (s)"
  )
```



Figure 3: Mean response times by trial type with 95% confidence intervals

Creating publication-quality figures (2)

```
corr_val <- cor(df$congruent, df$incongruent)

df %>%
  ggplot(aes(x = congruent, y = incongruent)) +
    geom_smooth(method="lm", alpha=.2) +
    geom_point() +
    annotate(
      "text",
      x = 20, y = 13,
      label=paste("r =", round(corr_val,3)),
      size = 6) +
    theme_bw(base_size=16) +
    theme(
      aspect.ratio=1,
      panel.grid = element_blank()
    ) +
    labs(
      x = "Congruent RT (s)",
      y = "Incongruent RT (s)"
    )
```

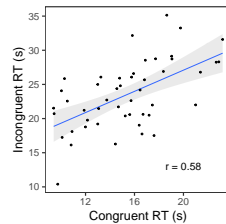


Figure 4: Scatterplot of Stroop task response times

Introduction to R

Plotting data

Creating publication-quality figures (2)

2021-08-24

Creating publication-quality figures (2)

```
corr_val <- cor(df$congruent, df$incongruent)

df %>%
  ggplot(aes(x = congruent, y = incongruent)) +
    geom_smooth(method="lm", alpha=.2) +
    geom_point() +
    annotate(
      "text",
      x = 20, y = 13,
      label=paste("r =", round(corr_val,3)),
      size = 6) +
    theme_bw(base_size=16) +
    theme(
      aspect.ratio=1,
      panel.grid = element_blank()
    ) +
    labs(
      x = "Congruent RT (s)",
      y = "Incongruent RT (s)"
    )
```



Figure 4: Scatterplot of Stroop task response times

Doing statistics

2021-08-24

Introduction to R
└─ Doing statistics

Doing statistics

```
t.test(df$congruent, df$incongruent, paired=TRUE)
```

```
##  
## Paired t-test  
##  
## data: df$congruent and df$incongruent  
## t = -13.681, df = 47, p-value < 2.2e-16  
## alternative hypothesis: true difference in means is not equal to 0  
## 95 percent confidence interval:  
## -9.324809 -6.933983  
## sample estimates:  
## mean of the differences  
## -8.129396
```

```
t.test(df$congruent, df$incongruent, paired=TRUE)  
  
##  
## Paired t-test  
##  
## data: df$congruent and df$incongruent  
## t = -13.681, df = 47, p-value < 2.2e-16  
## alternative hypothesis: true difference in means is not equal to 0  
## 95 percent confidence interval:  
## -9.324809 -6.933983  
## sample estimates:  
## mean of the differences  
## -8.129396
```

Basic linear regression

```
fit <- lm(incongruent ~ congruent, data=df)
summary(fit)
```

```
##
## Call:
## lm(formula = incongruent ~ congruent, data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -8.7229 -2.5923  0.1305  2.6973  9.1001
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   11.5641     2.5353   4.561 3.77e-05 ***
## congruent      0.7762     0.1607   4.831 1.55e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.076 on 46 degrees of freedom
## Multiple R-squared:  0.3365, Adjusted R-squared:  0.3221
## F-statistic: 23.33 on 1 and 46 DF, p-value: 1.549e-05
```

Introduction to R Doing statistics

Basic linear regression

```
fit <- lm(incongruent ~ congruent, data=df)
summary(fit)

##
## Call:
## lm(formula = incongruent ~ congruent, data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -8.7229 -2.5923  0.1305  2.6973  9.1001
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   11.5641     2.5353   4.561 3.77e-05 ***
## congruent      0.7762     0.1607   4.831 1.55e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.076 on 46 degrees of freedom
## Multiple R-squared:  0.3365, Adjusted R-squared:  0.3221
## F-statistic: 23.33 on 1 and 46 DF, p-value: 1.549e-05
```

2021-08-24

Hands-on exercise: Data in the wild

2021-08-24

Introduction to R
└ Hands-on exercise: Data in the wild

Hands-on exercise: Data in the wild

“How information about what is ‘healthy’ versus ‘unhealthy’ impacts children’s consumption of otherwise identical foods” (DeJesus et al., 2019).

In Experiment 1 of this study:

- Researchers presented children with two food options: one “healthy” and one “unhealthy”
- Then they left children alone with the food and allowed them to eat if they wanted
- A research assistant observed through a camera and recorded the number of bites the children took
- After each child finished, researchers weighed the remaining food in the dish to calculate how many grams of food the children ate

- Researchers presented children with two food options: one “healthy” and one “unhealthy”
- Then they left children alone with the food and allowed them to eat if they wanted
- A research assistant observed through a camera and recorded the number of bites the children took
- After each child finished, researchers weighed the remaining food in the dish to calculate how many grams of food the children ate

Load data from this paper yourself

The data is stored in the file `dejesus-example.csv`.

Data processing goals

Let's create a new tibble called `dejesus_cleaned` where we

- Include only the second trials for each child from Experiment 1
- Rename variables to be a bit clearer
- Remove some unnecessary variables

2021-08-24

Introduction to R

└─ Hands-on exercise: Data in the wild

└─ Hands-on: Loading the data

Hands-on: Loading the data

Load data from this paper yourself
The data is stored in the file `dejesus-example.csv`.

Data processing goals

Let's create a new tibble called `dejesus_cleaned` where we

- Include only the second trials for each child from Experiment 1
- Rename variables to be a bit clearer
- Remove some unnecessary variables

Variables we'd like

- **child_id**: participant number for each child
- **condition**: whether the child was told the food was healthy or unhealthy
- **bites**: how many bites the child took
- **grams**: how much of the food they ate (g)
- **age**: age of the children in years
- **gender**: gender of the children

Instructions

1. Look at the data with `head()`, `glimpse()` or `View()`
2. Filter for the 2nd trial from Exp. 1 using `filter()` (Exp. 1 is called "1_healthy_unh")
3. Rename variables to be a bit clearer using `rename()`
4. Remove unnecessary variables with `select()`

I didn't show you `select()`, take a look at the help with `?select` to learn about it.

Variables we'd like

- **child_id**: participant number for each child
- **condition**: whether the child was told the food was healthy or unhealthy
- **bites**: how many bites the child took
- **grams**: how much of the food they ate (g)
- **age**: age of the children in years
- **gender**: gender of the children

Instructions

1. Look at the data with `head()`, `glimpse()` or `View()`
2. Filter for the 2nd trial from Exp. 1 using `filter()` (Exp. 1 is called "1_healthy_unh")
3. Rename variables to be a bit clearer using `rename()`
4. Remove unnecessary variables with `select()`

I didn't show you `select()`, take a look at the help with `?select` to learn about it.

- why bother renaming things? did you know what `food_gen` meant?
- presumably the authors do, but it could be hard to remember

Now use the tools you've just learned to do some data sleuthing.

Ask yourself

- Are we missing data for any observations?
- Do all the observed values make sense?
- Can we fix or clean up these data? (how?)

Remember

- `head()` and `View()`
- `is.na()` and `any()`
- `filter()`
- plotting with `geom_point()`

Want them to: - check for NA - notice impossible values - maybe fix them - make a plot - do a test

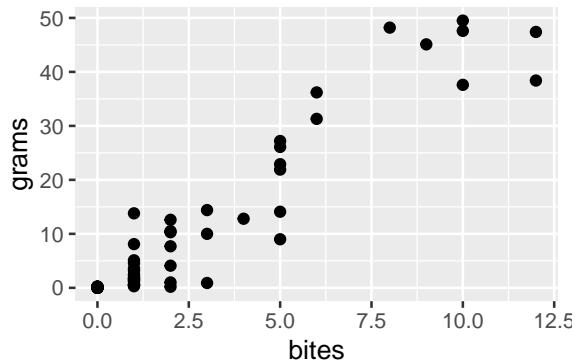
Now use the tools you've just learned to do some data sleuthing.

Ask yourself	Remember
• Are we missing data for any observations?	• <code>head()</code> and <code>View()</code>
• Do all the observed values make sense?	• <code>is.na()</code> and <code>any()</code>
• Can we fix or clean up these data? (how?)	• <code>filter()</code>
	• plotting with <code>geom_point()</code>

Hands-on: Inspecting with plotting

Plotting can help us look for outliers or possible data-entry errors.

```
dejesus %>%  
  ggplot(aes(x=bites, y = grams)) +  
  geom_point()
```

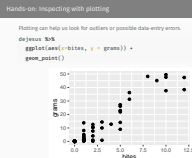


Introduction to R

Hands-on exercise: Data in the wild

Hands-on: Inspecting with plotting

2021-08-24



- A number of trials where children took zero bites show non-zero amounts of food being eaten according to grams
- These are most likely measurement errors, due to the inaccuracies of scales for small weights.
- If we like, we can fix these measurement errors using the `if_else()` function. The function takes 3 arguments, like so:

`if_else(condition, value_if_true, value_if_false)`

```
dejesus_cleaned <- dejesus %>%  
  mutate(  
    grams = if_else(bites==0, 0, grams)  
  )
```

- scale is supposed to have .1g accuracy, which means 1sd to measurements. .1 to .2 is within 1-2sd. if we think a .5g bite is possible, then .3g observation could be 2sd off a .5g bite, rather than 3sd off a 0g bite. But it's hard to know without knowing the accuracy of the "bite" observers.
- We could treat this as a statistical inference problem, but lets save ourselves for more interesting issues. We can probably safely recode the .1 and .2 on zero bites as zeros. It's not as clear what to do with the .3 on 1 bite.

```
• A number of trials where children took zero bites show non-zero amounts of  
food being eaten according to grams  
• These are most likely measurement errors, due to the inaccuracies of scales for  
small weights.  
• If we like, we can fix these measurement errors using the if_else() function.  
The function takes 3 arguments, like so:  
if_else(condition, value_if_true, value_if_false)  
dejesus_cleaned <- dejesus %>%  
  mutate(  
    grams = if_else(bites==0, 0, grams)  
  )
```

- Calculate the gender and age breakdown of the children

```
summary(dejesus$age)
```

```
dejesus %>% count(gender)
```

- Plot the bites and/or grams eaten data by condition
- Try a few different approaches: `geom_boxplot()`, `geom_violin()`, `geom_jitter()`
- Try plotting mean with standard error bars

```
ggplot(dejesus, aes(x=grams)) + geom_histogram()  
ggplot(dejesus) %>% aes(x=condition, y = bites) +  
geom_jitter(width=.2)
```

- Use `lm()` to fit a simple linear regression model predicting grams eaten from number of bites the children took.
- From the output of the regression, what is the expected size of a bite in grams?
- Does the model's intercept make sense?

- I already showed you a t-test, so try using a non-parametric test with `wilcox.test()` to compare the groups in terms of grams or bites

2021-08-24

Introduction to R

└ Hands-on exercise: Data in the wild

└ Hands-on: A statistical test

• I already showed you a t-test, so try using a non-parametric test with `wilcox.test()` to compare the groups in terms of grams or bites

- You have now seen all the basic workflow steps for doing statistics with R
 - loading data
 - processing data
 - plotting data
 - making models and performing tests

- discuss homework submission

Extra slides

2021-08-24

Introduction to R
└ Extra slides

Extra slides

Modifying categorical variables

The three commands below make the same changes.

```
df_long %>%
  mutate(trial_type = ifelse(trial_type=="congruent", "Congruent", "Incongruent"))

df_long %>%
  mutate(
    trial_type = fct_recode(
      trial_type,
      "Congruent"="congruent",
      "Incongruent"="incongruent")
  )

df_long %>%
  mutate(
    trial_type = case_when(
      trial_type=="congruent" ~ "Congruent",
      trial_type=="incongruent" ~ "Incongruent"
    )
  )
```

Introduction to R

Extra slides

Modifying categorical variables

Modifying categorical variables

The three commands below make the same changes.

```
df_long %>%
  mutate(trial_type = ifelse(trial_type=="congruent", "Congruent", "Incongruent"))

df_long %>%
  mutate(
    trial_type = fct_recode(
      trial_type,
      "Congruent"="congruent",
      "Incongruent"="incongruent")
  )

df_long %>%
  mutate(
    trial_type = case_when(
      trial_type=="congruent" ~ "Congruent",
      trial_type=="incongruent" ~ "Incongruent"
    )
  )
```

2021-08-24